



Welcome to the Melbourne OpenShift Online Meetup



OPENSIFT



ANSIBLE

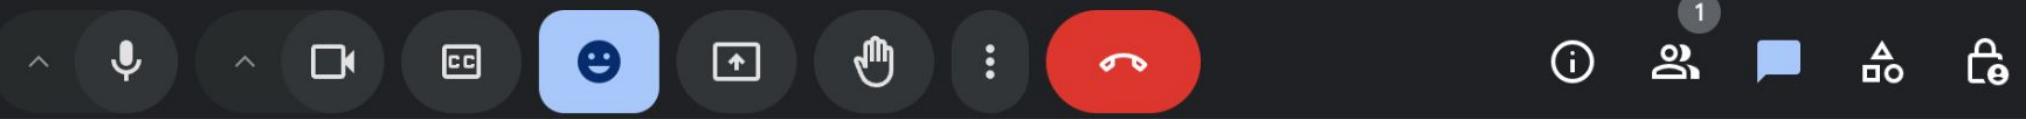


Red Hat
Enterprise Linux AI



Red Hat
OpenShift AI

Send a message to everyone





**Red Hat
Summit**

Agenda:

- ▶ Red Hat Summit Updates
- ▶ Batch Jobs in OpenShift
- ▶ Product Roadmaps

Red Hat
Summit

AnsibleFest

Boston, MA | May 19-22, 2025

RH Summit Updates



RHEL

Red Hat Enterprise Linux 10 addresses these market forces



Tight resource
constraints



Accelerated
cloud adoption



Security
threats



Artificial
Intelligence



Red Hat Enterprise Linux 10 will help you...



- ⌚ **Address the Linux skills gap**
with decades of Red Hat's Linux knowledge and expertise
- ✂️ **Contain drift and accelerate delivery**
with container tools and technologies
- ⌚ **Make better decisions at build time**
when it's typically easier and cheaper to make changes
- ⚠️ **Resist security attacks from hackers**
when quantum computers become prevalent
- 💡 **Leverage Red Hat Enterprise Linux as a trusted AI foundation**
with an extensive ecosystem of trusted partners and tools



Red Hat Enterprise Linux 10 will help you...

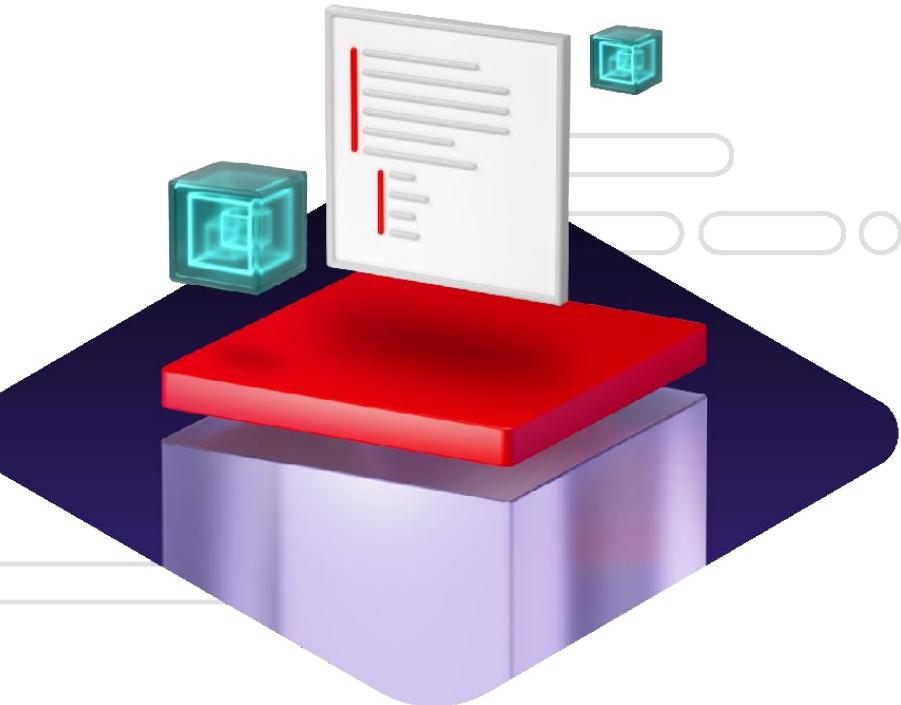


- ⌚ Address the Linux skills gap
with Command Line Assistant, powered by RHEL Lightspeed
- 🕒 Contain drift and accelerate delivery
with image mode for Red Hat Enterprise Linux
- 🕒 Make better decisions at build time
with Insights lifecycle planning and recommendations
- ⚠ Resist security attacks from hackers
with new quantum-resistant algorithms (and more to come)
- 💡 Leverage Red Hat Enterprise Linux as a trusted AI foundation
with the RHEL extensions repository, Podman Desktop, and the AI partner validation program



Contain drift and accelerate delivery

using container tools and technologies



With image mode for Red Hat Enterprise Linux,
you can:

- ▶ **Speed time to market**
using DevOps and CI/CD practices, which now include the OS
- ▶ **Streamline operations**
by automating updates and rollbacks—just like your smartphone
- ▶ **Enhance security**
by reducing your attack surface with immutable system images
- ▶ **Simplify appliance creation**
by combining the OS with apps and drivers for faster development and delivery

Because systems should be as easy to update as smartphones



Standardizing and innovating with containers

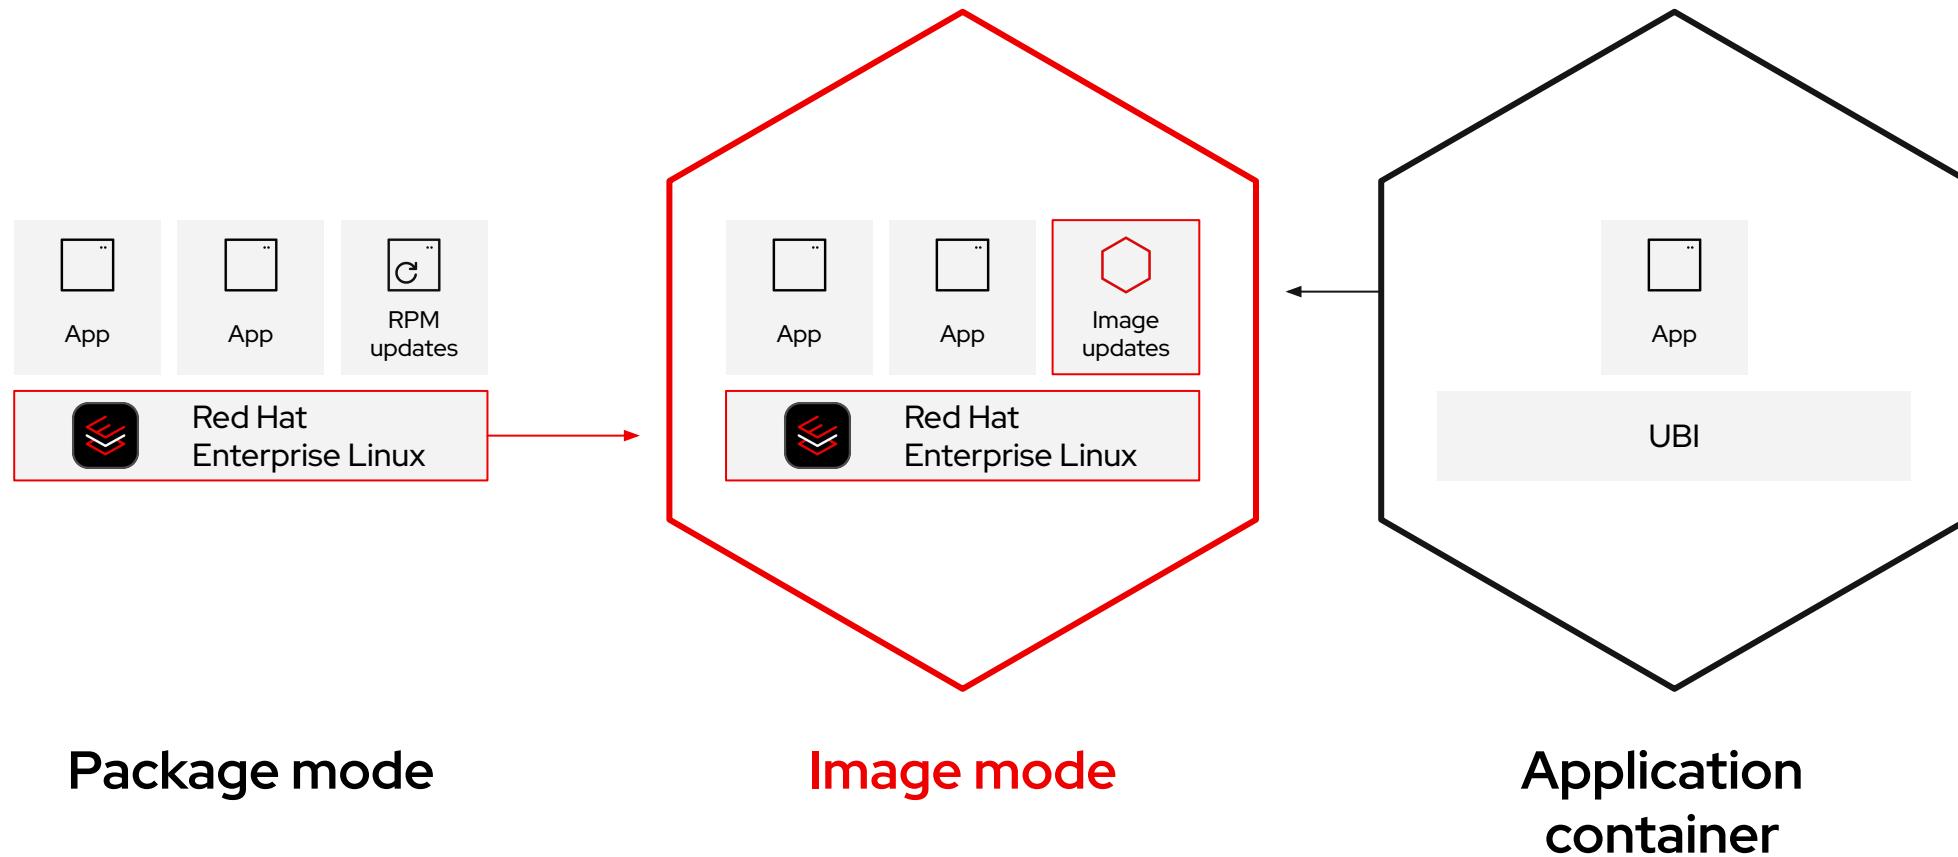


Image mode for Red Hat Enterprise Linux

A container-native workflow for the life cycle of a system

```
FROM rhel10/rhel-bootc:latest

RUN dnf install -y [software]
[dependencies] && dnf clean
all

ADD [application]
ADD [configuration files]

RUN [config scripts]
```

Build

Define your entire system—OS, applications, and dependencies—with just a bootc base image and container file. Leverage your existing container tools and pipelines for rapid image creation and testing.

Deploy

Easily convert to VM/cloud images, deploy on bare metal via the Red Hat Enterprise Linux installer, or even reinstall on existing cloud images using bootc.

Manage

Engineered for modern GitOps and CI/CD workflows. Fully drive and automate systems via pipelines or scale control through Red Hat Insights, Satellite, and Ansible.



Resist security attacks from hackers

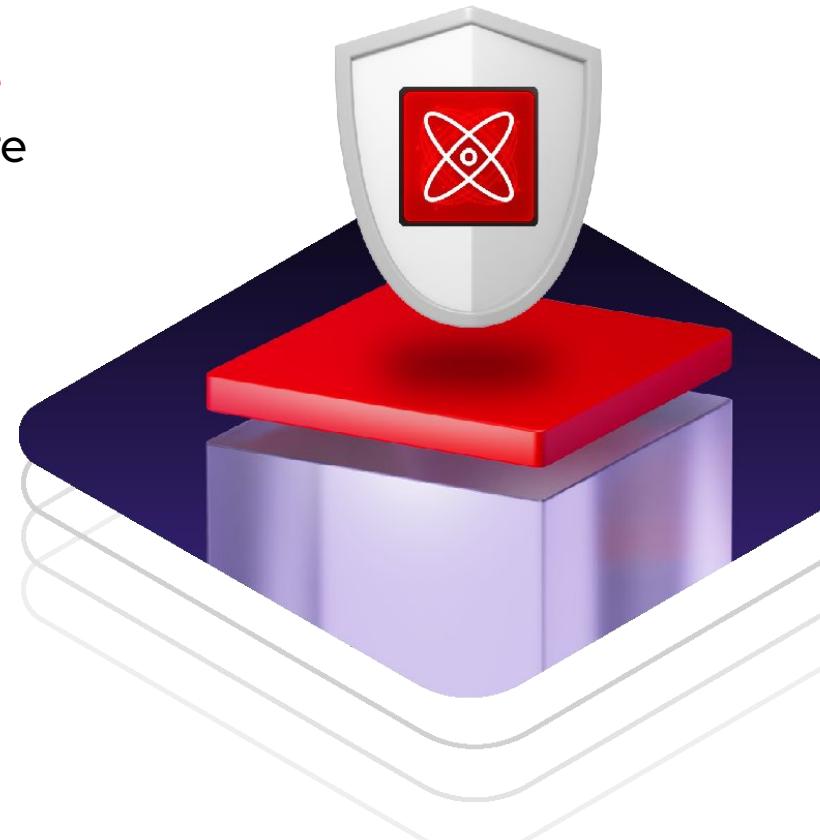
when quantum computers become prevalent

- ▶ Red Hat Enterprise Linux 10 is the first enterprise Linux distribution to be **post-quantum capable** with new quantum-resistant algorithms (and more to come) so you can prepare now for future compliance mandates
 - OpenSSL
 - ML-KEM (FIPS 203)
 - ML-DSA (FIPS 204)
- ▶ A new FIPS module

simplifies development of FIPS-compliant app deployments on containers, VMs or bare metal

- ▶ FIPS cryptographic standards can be validated separately

meaning CVE fixes related to Open SSL or crypto libraries can be done *without* requiring a new FIPS validation certificate (average time to obtain is >300 days!)



The best defense is a strong partner.



Addressing the Linux skills gap

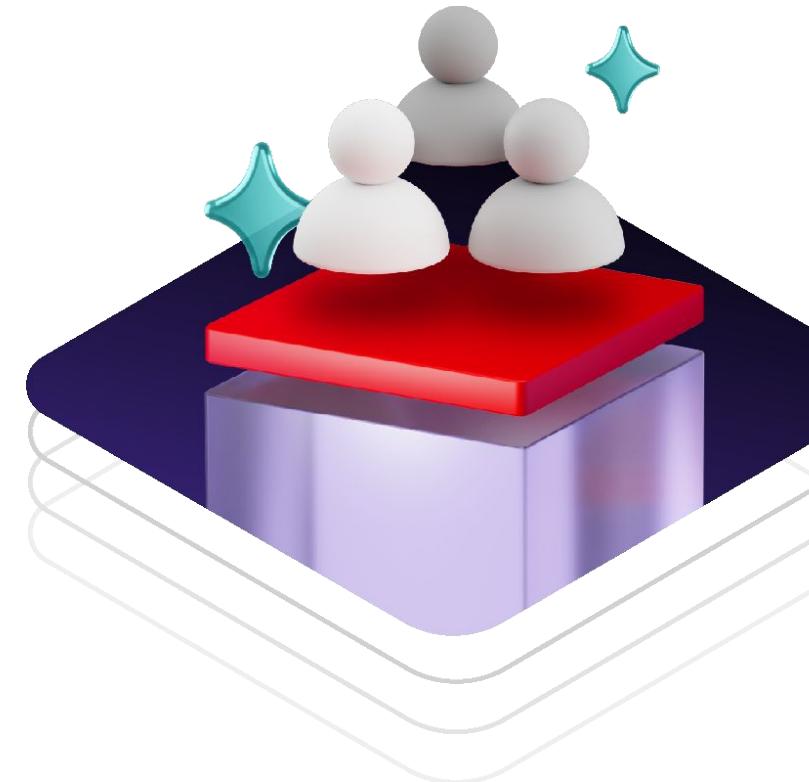
with decades of Red Hat's Linux knowledge and expertise

Red Hat Enterprise Linux Lightspeed

combines decades of Red Hat Enterprise Linux expertise with AI technologies to proactively inform and simplify how both newer and experienced IT professionals build, deploy, and manage Red Hat Enterprise Linux.

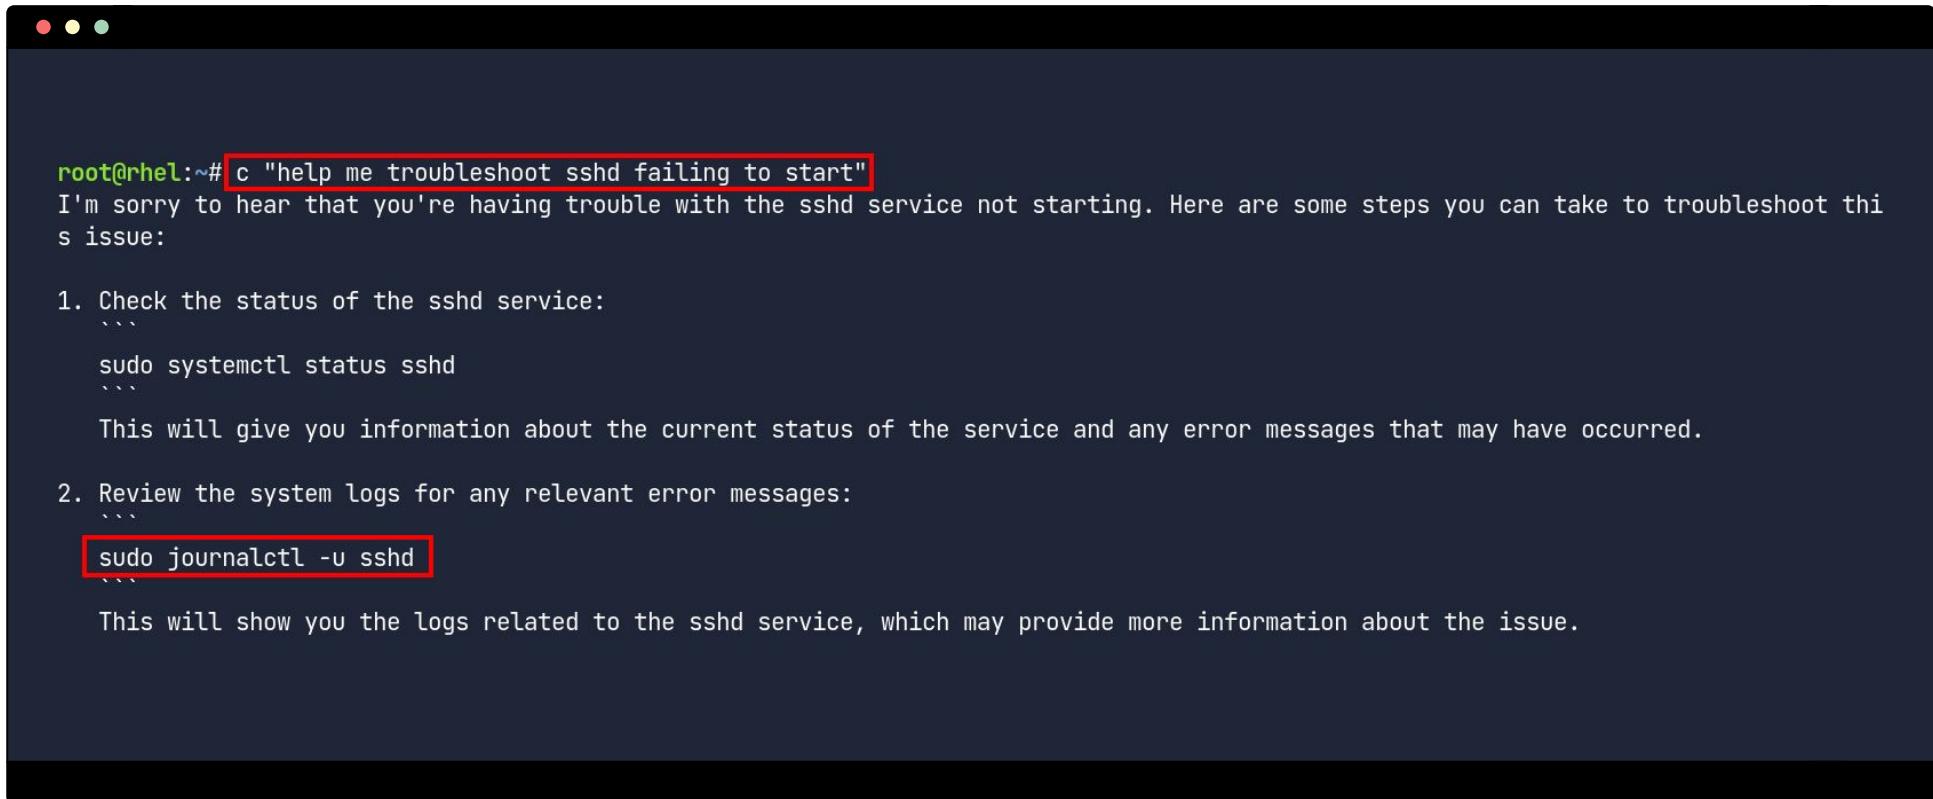
- ▶ Use plain language to simplify the way you interact with Red Hat Enterprise Linux
- ▶ Make better decisions with recommendations and actionable guidance

Simplify tasks. Amplify results.



A new command line assistant

Powered by Red Hat Enterprise Linux Lightspeed

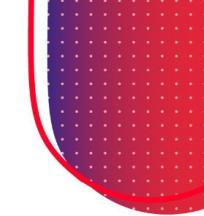


The screenshot shows a terminal window with a dark background. At the top, there are three small colored dots (red, yellow, green) representing window control buttons. The terminal prompt is "root@rhel:~#". Below the prompt, a command is entered: "c \"help me troubleshoot sshd failing to start\"". This command is highlighted with a red rectangular box. The terminal then displays a troubleshooting message: "I'm sorry to hear that you're having trouble with the sshd service not starting. Here are some steps you can take to troubleshoot this issue:". The first step is listed as "1. Check the status of the sshd service:" followed by three ellipses ("..."). Below this, the command "sudo systemctl status sshd" is shown, also preceded by three ellipses. A note states: "This will give you information about the current status of the service and any error messages that may have occurred.". The second step is listed as "2. Review the system logs for any relevant error messages:" followed by three ellipses ("..."). Below this, the command "sudo journalctl -u sshd" is shown, also preceded by three ellipses. A note states: "This will show you the logs related to the sshd service, which may provide more information about the issue."

```
root@rhel:~# c "help me troubleshoot sshd failing to start"
I'm sorry to hear that you're having trouble with the sshd service not starting. Here are some steps you can take to troubleshoot this issue:
1. Check the status of the sshd service:
...
sudo systemctl status sshd
...
This will give you information about the current status of the service and any error messages that may have occurred.
2. Review the system logs for any relevant error messages:
...
sudo journalctl -u sshd
...
This will show you the logs related to the sshd service, which may provide more information about the issue.
```



AAP



AAP Highlights from Summit

Collections

- New Partner Collections

Policy Enforcement with AAP

- Introduction to [Policy enforcement interactive demo](#)
- [Use case Automated Policy as Code with Red Hat Ansible Automation Platform](#)

Ansible Lightspeed

- Docs: [Deploying the Ansible Lightspeed intelligent assistant on OpenShift Container Platform](#)
- Interactive experience: [Test drive the Ansible Lightspeed intelligent assistant](#)

AIOps

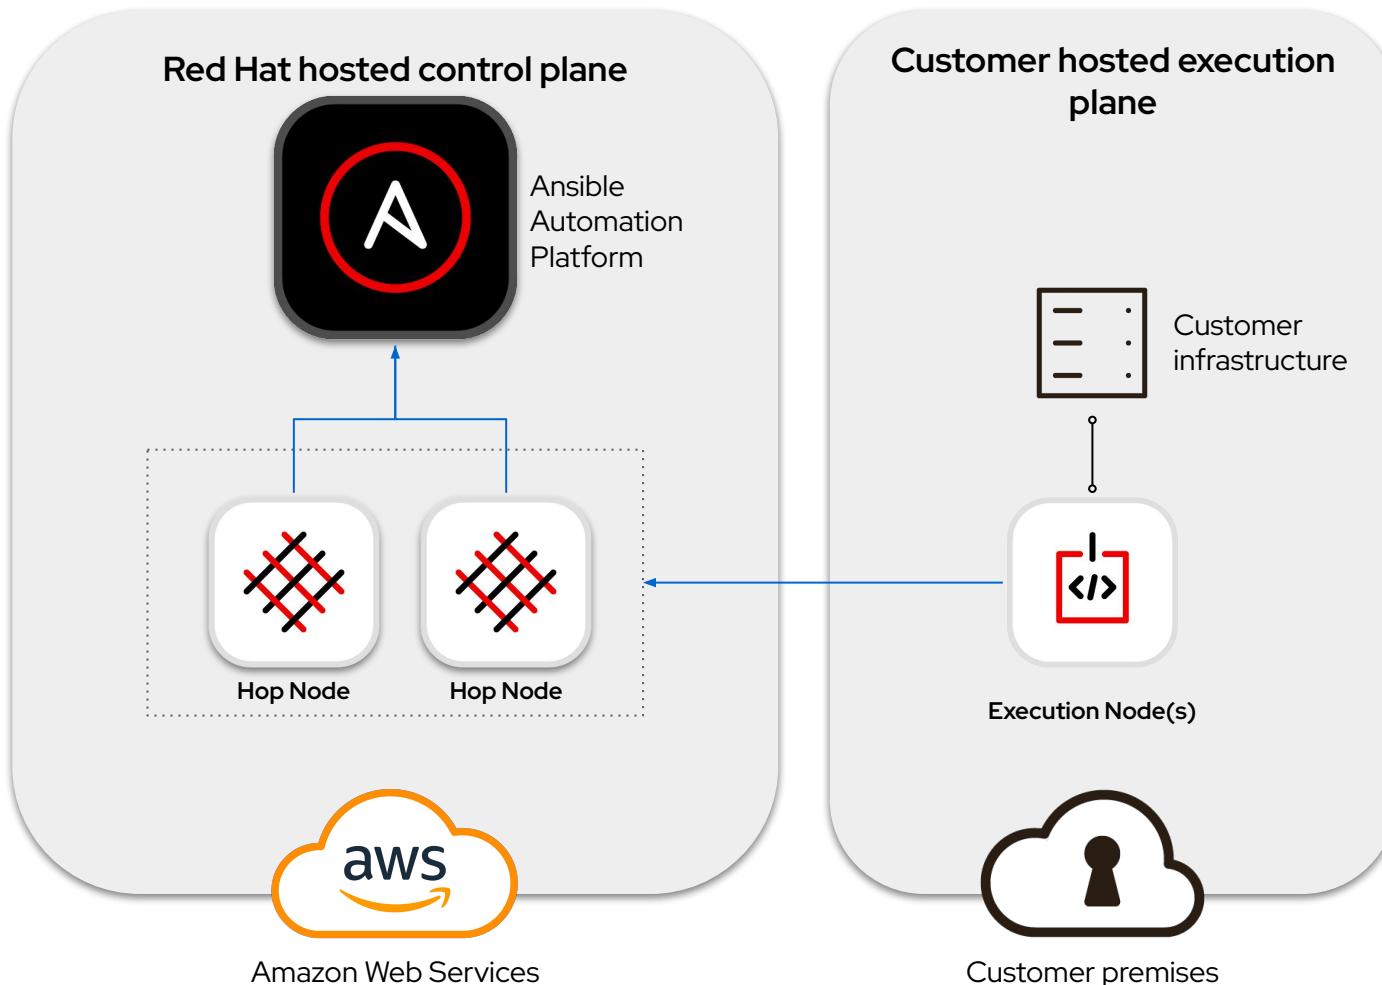
- [AIOps automation with Ansible – Solution Guide](#)

Self Service

- Aimed at Domain SME (non-technical) users unfamiliar with AAP/automation.
- [AAP Self-service automation \(Technology Preview\)](#)

Ansible Automation Platform Service on AWS

Focus on innovation through automation instead of managing the platform



- ▶ **Red Hat Ansible Automation Platform Service on AWS** is a fully managed Ansible Automation Platform control plane
- ▶ **Execution nodes** - deployed in your infrastructure, with the ability to pull automation tasks vs push, simplifying networking and minimizing network security configuration
- ▶ **Purchase in the AWS marketplace** and draw down against committed spend agreements

<https://www.redhat.com/en/blog/red-hat-ansible-automation-platform-service-aws-now-available-aws-marketplace>

Policy enforcement

Hashicorp
Terraform and
Vault

Automation
Dashboard

Cloud Marketplace
offerings

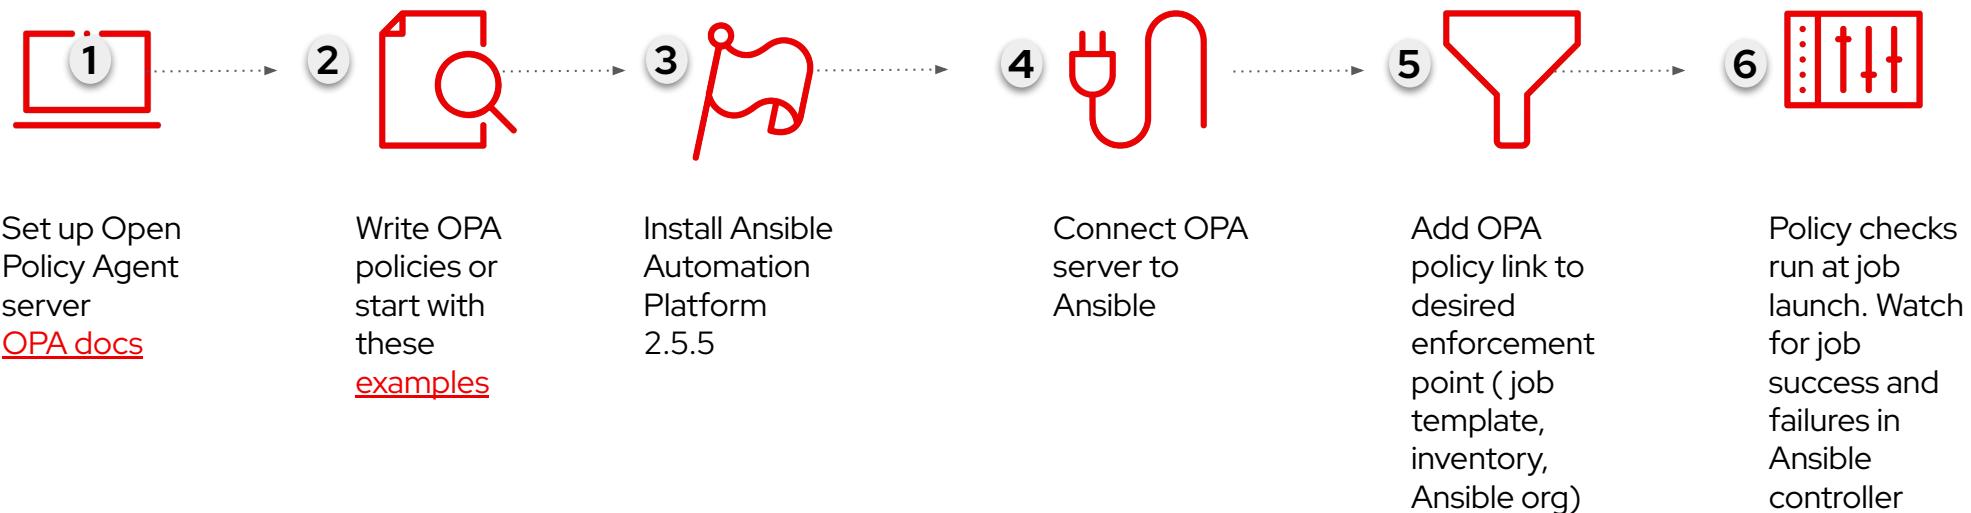
Node counting

Self-service
automation

2.5 upgrade
updates

Getting started with policy enforcement

Requirements & Configuration



AI

Key Announcements



New product announcement.
Optimizes model inference creating faster and more cost-effective model deployments



Validated Models



Expanded capabilities and offering channels

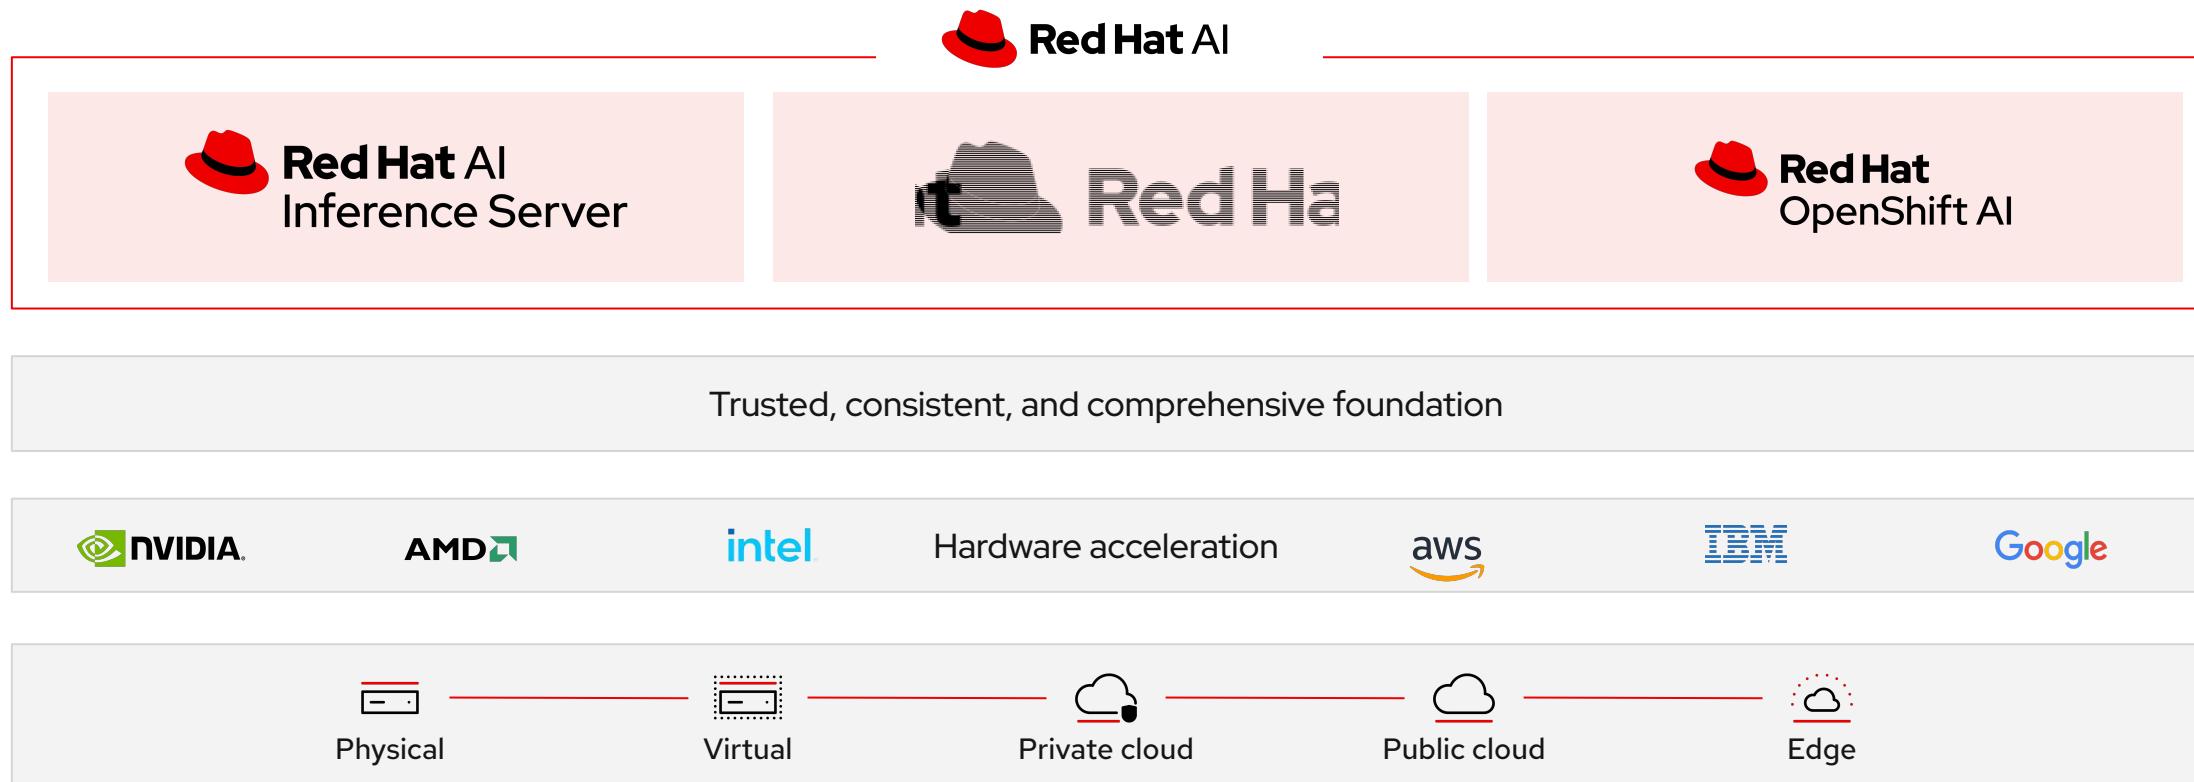


Expanded capabilities





Unified Platform for Generative and Predictive AI





Unified Platform for Generative and Predictive AI



**Red Hat AI
Inference Server**



**Red Hat
OpenShift AI**

**For customers who need Gen AI
model Inference on RHEL/Linux or
OpenShift/Kube**



Physical

Virtual

Private cloud

Public cloud



Edge

Red Hat AI Inference Server

In 78 words or less...

GPU's enable GenAI workloads. IDC predicts Inferencing is expected to be 80% of GPU usage (compare to say training) in coming years.

Red Hat AI Inference Server optimizes model inference across the hybrid cloud and edge, creating **faster and more cost-effective model deployments**.

It allows customers to **run any generative AI model on any hardware accelerator** (NVIDIA, Intel, AMD) and **in any environment** (datacenter, cloud, edge)—providing flexibility and choice to meet business requirements.

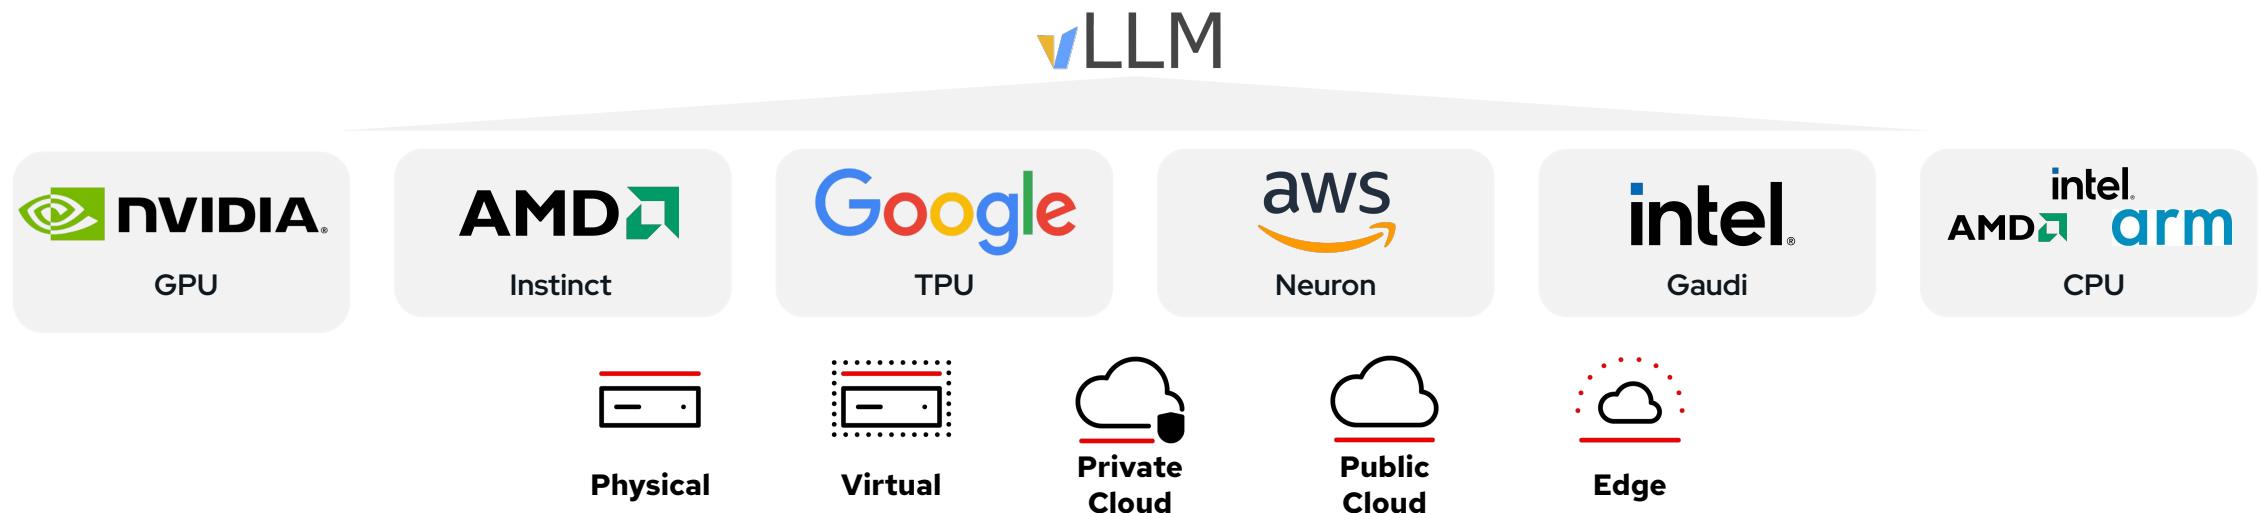
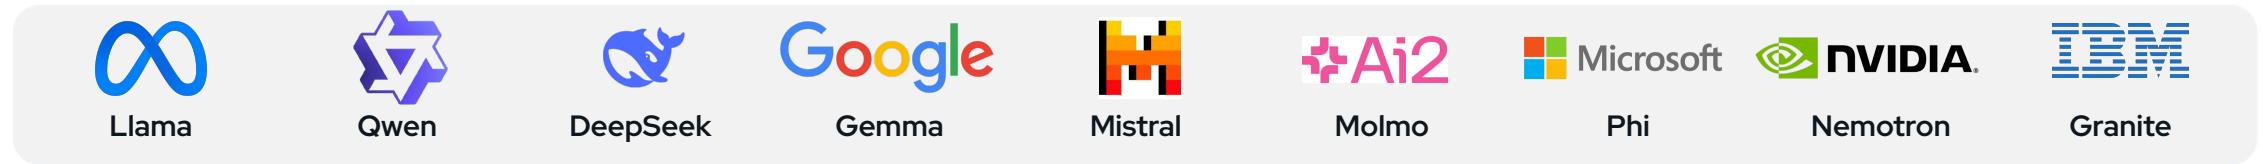
Its runtime, vLLM: **maximizes throughput, minimises latency** and **optimises utilisation and cost**.

An optimized model repository **accelerates model serving**, while an LLM compressor helps **reduce GPU utilization while preserving accuracy**.



Cross Platform

vLLM is emerging as the Linux of GenAI Inference



Single Platform To Run Your Models Across Accelerators and OEMs.

Red Hat AI repository on Hugging Face

A collection of third-party validated and optimized large language models

Broad Collection of models



Llama



Qwen



Gemma



Mistral



DeepSeek



Phi



Molmo

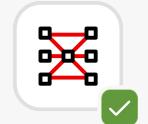


Granite



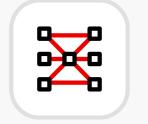
Nemotron

Validated models



- ▶ Tested using realistic scenarios
- ▶ Assessed for performance across a range of hardware
- ▶ Done using GuideLLM benchmarking and LM Eval Harness

Optimized models



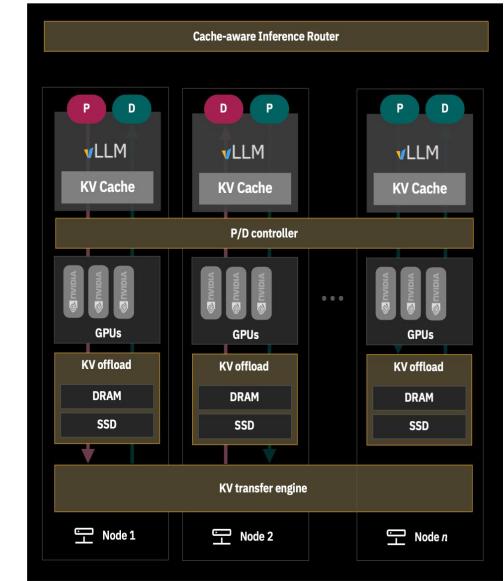
- ▶ Compressed for speed and efficiency
- ▶ Designed to run faster, use fewer resources, maintain accuracy
- ▶ Done using LLM Compressor with latest algorithms

New open source project: llm-d

Distributed inference at scale

Flexible and Distributed architecture to meet SLOs efficiently

- **Consumers** demand low latency and high accuracy
- Diversity of **Models and sizes, context lengths**
- Limited availability and access to **variety of AI Accelerators**
- **System SLOs:** Support N concurrent connections, with Y latency, and Z accuracy
- Meet **Token Cost** targets (Cost per 1 Million Tokens)

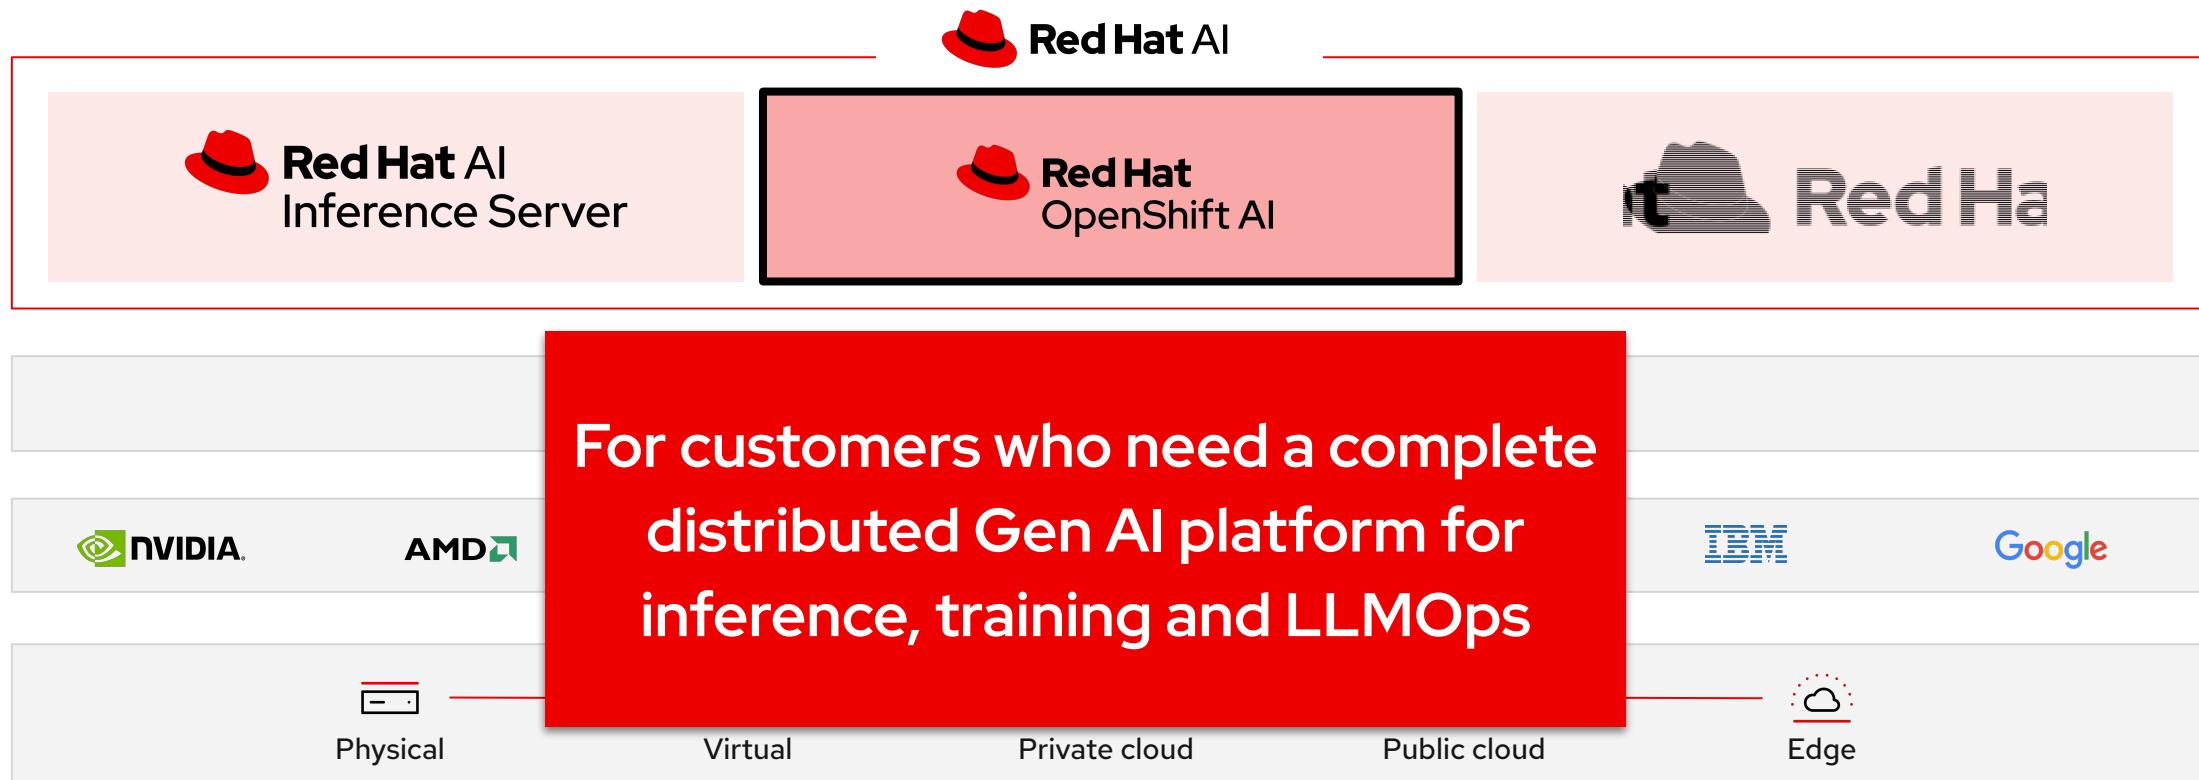


Benefits

- Maximize Token Revenue
- Reduce Token Costs
- Boost Latency and Throughput
- Seamless Scaling



Unified Platform for Generative and Predictive AI



Red Hat OpenShift AI

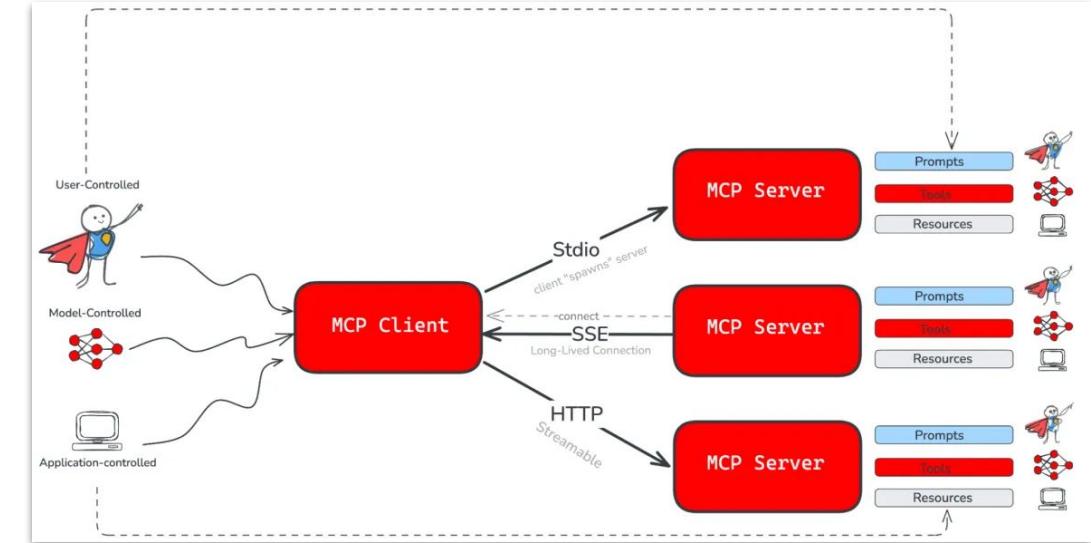
Agentic AI

A flexible foundation for AI agent development

- Streamlined development via an AI API server
- Direct access to core components

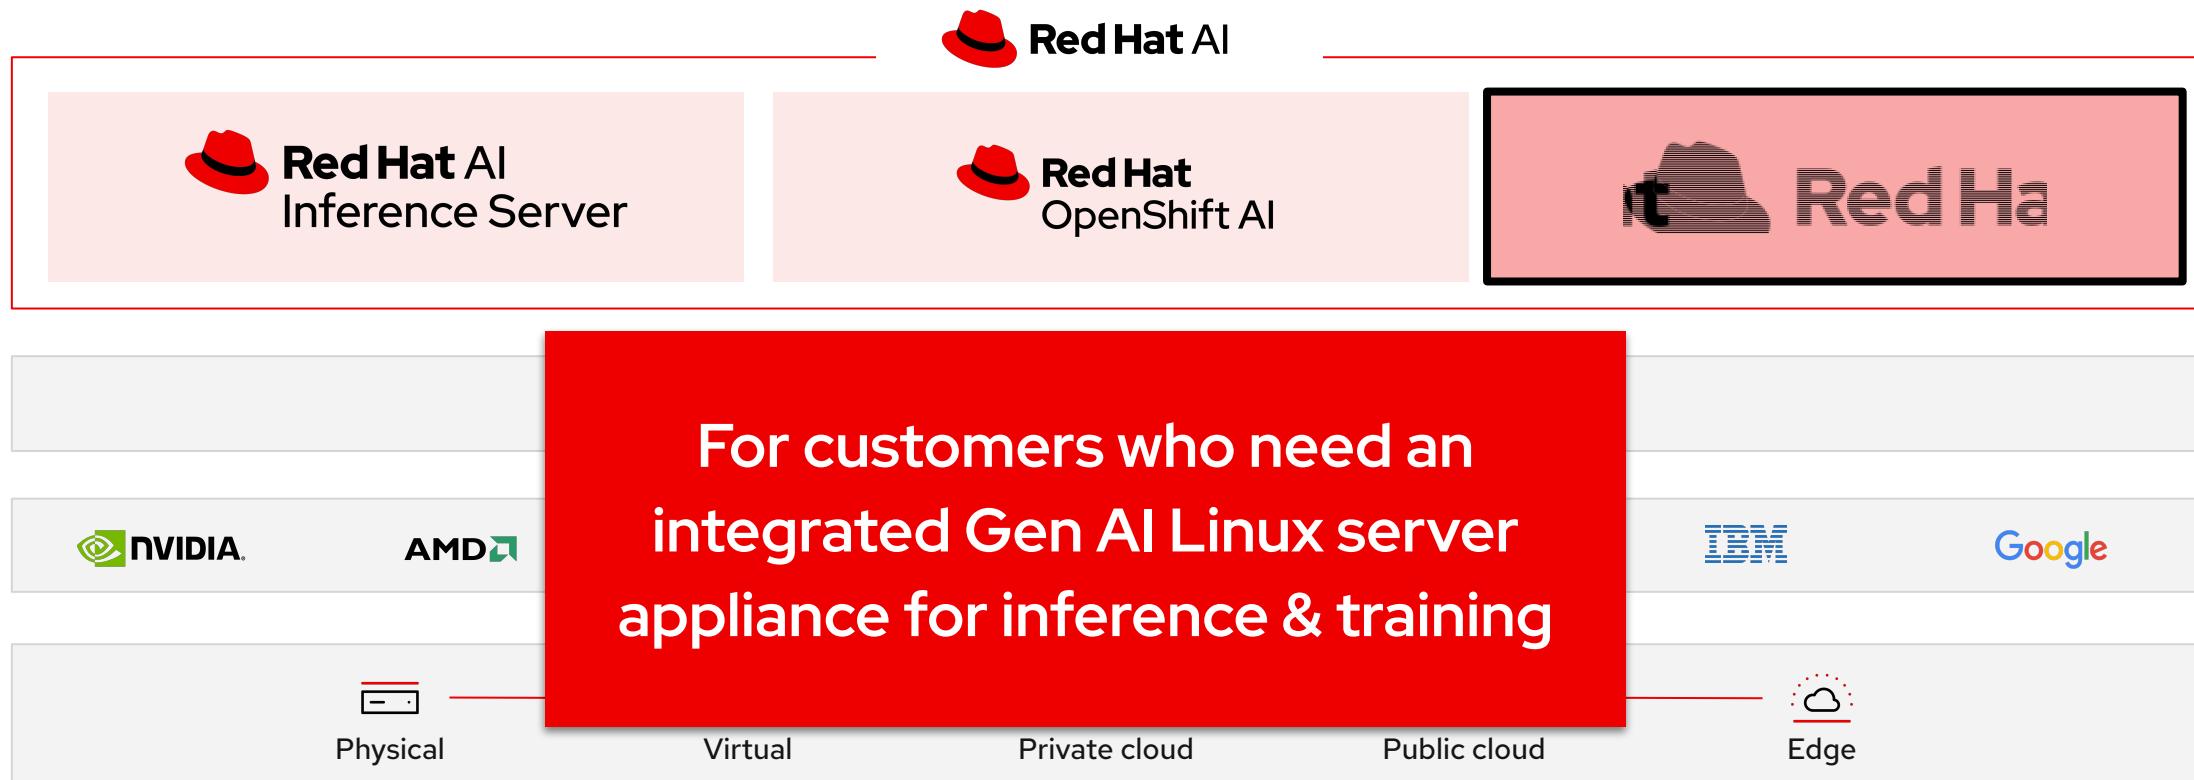
A unified AI API server on OpenShift AI with Llama Stack

- **Standard AI APIs:** Llama Stack offers consistent endpoints for core AI tasks (inference, RAG, agents, safety).
- **Abstracts complexity:** This allows AI developers to switch tools or backends without rewriting significant portions of their code.
- **Extensible design:** Its provider model allows plugging in different backends (e.g., vector databases) via standard interfaces, similar to Kubernetes' CRDs and operators.
- **Agent-focused:** Includes built-in concepts critical for agent development .





Unified Platform for Generative and Predictive AI



Red Hat's AI Building Blocks

User/Benefit:

Immediate Insights

Private Data Analysis

Legally Bounded via Guardrails

AI-Driven Rapid Innovation

Applications:

Chatbot

Speech-to-Text

Summariser

Object Detection

AI-Driven Observability

AI-Driven Automation

Reduced barrier to entry

Tooling:



High Productivity

Platforms:



LLM's - Foundational, Extended, Custom



Guardrails

LLM Community



Granite Model

Validated models

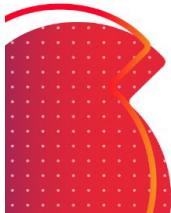
1000's of other models



Trusty AI
Provenance and Governance

- Governance - where is your model from?
- Was it trained on copyright data?
- Was it trained on toxic data?
- Is it using Scarlett Johansson's voice?

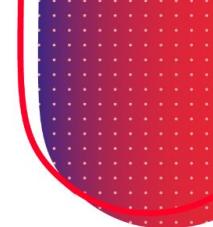
Best of both worlds: smaller taylorable models



Red Hat
Summit

AnsibleFest

Boston, MA | May 19-22, 2025



Summary Highlights from Summit

AI

- "Any Model, Any Accelerator, Any Cloud"
- **Red Hat Inference Server** based on [vLLM](#)
- New Open source project [llm-d](#) as "Kubernetes of model serving"
- Red Hat AI Validated Models
- MCP Servers for RH Products Coming Real Soon

OpenShift Lightspeed

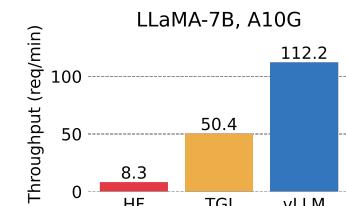
- (June) GA of Red Hat OpenShift Lightspeed, a generative AI-based virtual assistant integrated into Red Hat OpenShift.
- [Interactive Walkthrough: OpenShift Lightspeed](#)

RHEL

- RHEL Image Mode
- Try [Lab 1 - Introduction to image mode for Red Hat Enterprise Linux](#)
- Try [Lab 2 Day 2 operations with image mode for Red Hat Enterprise Linux](#)
- Post-Quantum Cryptography
- RHEL Lightspeed
- RHEL Security Select



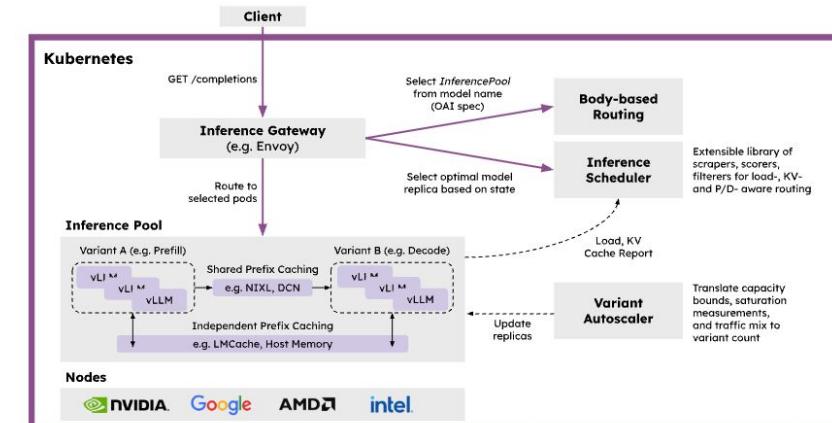
Easy, fast, and cheap LLM serving for everyone



With vLLM, LMSYS was able to cut the number of GPUs used for serving the above traffic by 50%.

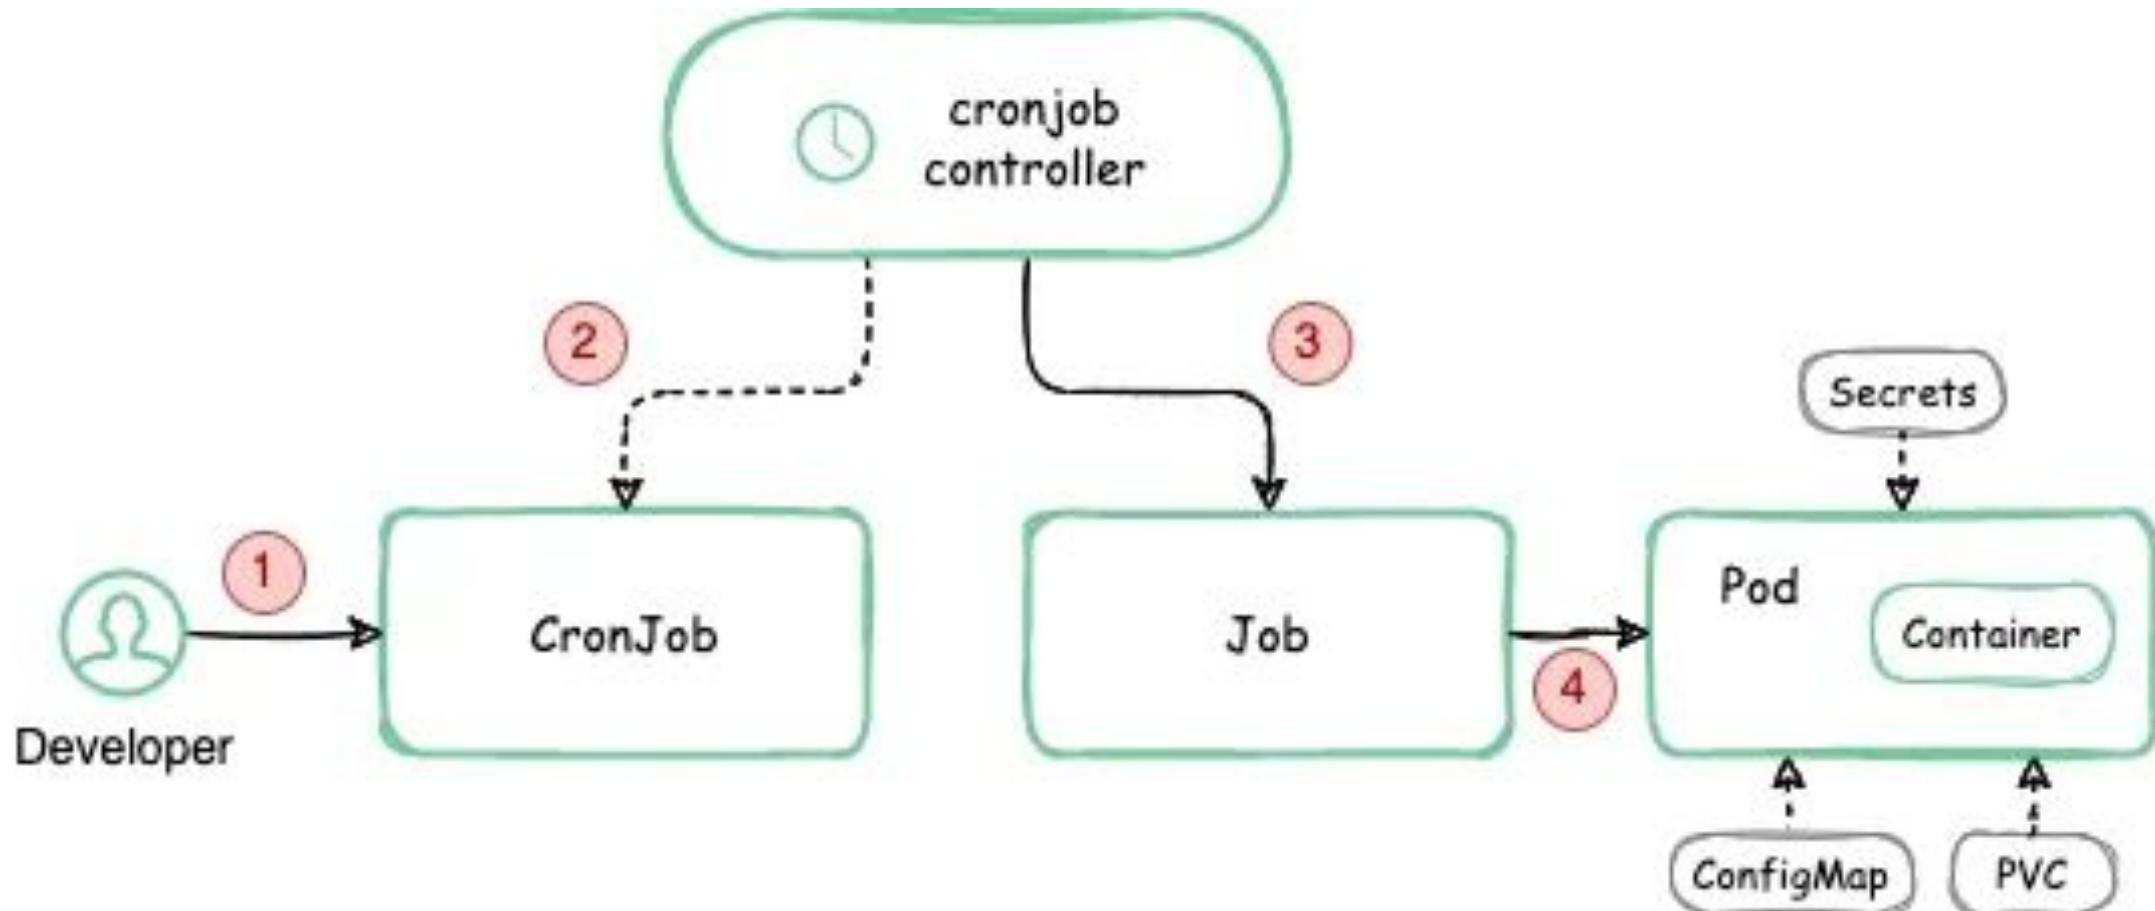


A Kubernetes-native high-performance distributed LLM inference framework
<https://llm-d.ai/>





Batch Jobs in OpenShift



When to Use CronJobs

- **Regularly Scheduled Tasks:** CronJobs are best suited for jobs that need to run on a regular, predictable schedule. This can range from system maintenance tasks like cleanup and backup to data processing tasks that need to run during low traffic times.
- **Polling:** If you need to periodically check an external service for updates, CronJobs can be a good fit. *Example:* In a microservices architecture, if a service needs to verify the status of a remote API endpoint periodically, a CronJob can be set to hit that endpoint at fixed intervals.
- **Notifications and Reminders:** If you need to send notifications or reminders at certain times of the day or week, you can use CronJobs to schedule these tasks. *Example:* If you have a service that sends weekly newsletters to your users, a CronJob can be set up to run this service every week at a specified time.

When Not to Use CronJobs

- **Short, Frequent Jobs:** Kubernetes has some overhead in scheduling and running jobs, so if you have a job that needs to run very frequently (like every few seconds), a CronJob might not be the best choice. A continuously running service that manages its internal timing could be more suitable in this case.
- **Non-periodic, Trigger-based Jobs:** If your jobs need to run based on certain events or triggers (e.g., a new file appearing in a directory, or a message arriving on a queue), CronJobs may not be the best fit. Instead, event-driven architectures or workflows should be used.
- **Long Running or Continuous Jobs:** CronJobs are intended for jobs that run to completion. For long-running or continuous tasks that don't have a clear end, such as a web server, CronJobs are not suitable.
- **Jobs with Complex Dependencies:** If you have jobs with complex dependencies, where one job must complete before another begins, and these relationships can't easily be expressed with the Cron syntax, you might need a more sophisticated workflow management system.

Workloads ▾

Project: batchjobs ▾

CronJobs

Name	Schedule	Suspend	Concurrency policy
CJ example	@daily	False	Allow
CJ example2	@daily	False	Allow
CJ example4	*****	False	Replace

Example Simple Cron Job

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: load-generator-${X}
spec:
  schedule: "* * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: quick-success
              image: image-registry.openshift-image-registry.svc:5000/openshift/tools:latest
              imagePullPolicy: IfNotPresent
              command:
                - echo
                - success
      restartPolicy: OnFailure
```

Example Indexed Job

```
kind: CronJob
apiVersion: batch/v1
metadata:
  name: example4
  namespace: batchjobs
spec:
  schedule: '* * * * *'
  concurrencyPolicy: Replace
  suspend: false
  jobTemplate:
    metadata:
      creationTimestamp: null
    spec:
      parallelism: 3
      completions: 3
      template:
        metadata:
          creationTimestamp: null
        spec:
          volumes:
            - name: input
              downwardAPI:
                items:
                  - path: data.txt
                    fieldRef:
                      apiVersion: v1
                      fieldPath: 'metadata.annotations["batch.kubernetes.io/job-completion-index"]'
        containers:
          - name: indexjob
            image: 'image-registry.openshift-image-registry.svc:5000/openshift/tools:latest'
            command:
              - bash
              - '-c'
              - |
                echo "My partition: ${JOB_COMPLETION_INDEX}"
                cat /input/data.txt
            volumeMounts:
              - name: input
                mountPath: /input
            terminationMessagePath: /dev/termination-log
            terminationMessagePolicy: File
            imagePullPolicy: Always
        restartPolicy: Never
        terminationGracePeriodSeconds: 30
        ttlSecondsAfterFinished: 100
        completionMode: Indexed
        suspend: false
        successfulJobsHistoryLimit: 1
        failedJobsHistoryLimit: 1
```

To create an Indexed Job, you just have to add `completionMode: Indexed` to the Job spec and make use of the `JOB_COMPLETION_INDEX` environment variable or mount it as shown

Use [Indexed Jobs](#) when you need to partition workload and run in parallel.

Project: batchjobs ▾

Jobs > Job details

J example4-29163208 SuccessCriteriaMet

Details YAML Pods Events

Filter Name Search by name... /

Name ↑	Status ↓
P example4-29163208-0-5rgjm	Completed
P example4-29163208-1-55wzq	Completed
P example4-29163208-2-b9fhj	Completed

My partition: 0
0

My partition: 1
1

My partition: 2
2



Product RoadMaps

AI & Middleware

Red Hat Application Foundations

- Application and API connectivity
 - Service Interconnect,
 - 3Scale
 - Connectivity Link
- Messaging
 - Streams for Apache Kafka
 - AMQ broker
- JBoss EAP
- Enterprise Integration & Data transformation
 - Red Hat build of Apache Camel 4

Developer Tooling & Onboarding

- [VSCode IDE Plugins](#)
- [Podman Desktop](#)
- [Install OpenShift on Desktop / Microshift](#)
- [Red Hat Developer Hub](#)
- [Validated & Solution Patterns](#)

Included Software with OpenShift

- RHEL and RHEL CoreOS.
- OpenShift Virtualization (includes RHEL hosts)
- Popular languages and runtimes
 - PHP, Python, Perl, Node.js, Ruby, OpenJDK, [.NET](#)
 - Red Hat build of **Quarkus**
- Databases, such as MySQL, Postgres..
- JBoss Web Server (Tomcat)
- Red Hat build of Keycloak / SSO
- Dev Spaces (IDE in a Container)
- Pipelines & GitOps
- Serverless
- Service Mesh

Red Hat AI

- [RHEL AI](#)
- [OpenShift AI](#)
- [Red Hat AI Inference Server](#)



OpenJDK

What's New

Here are the highlights of OpenJDK 21:

- ▶ UTF-8 by default
- ▶ Simple Web Server
- ▶ Code snippets in Java API documentation
- ▶ Virtual Threads
- ▶ Please refer to the [release notes](#) for OpenJDK 21
- ▶ OpenJDK 11 Extended Lifecycle Support began in October 2024
- ▶ Cryostat 4.0 (Full Support, [What's new](#))

Coming soon...

Some of the features included in our roadmap for the next 12 months:

- ▶ OpenJDK 25 in Fall 2025
- ▶ Critical Patch Updates (Every three months) for JDK:
 - 8, 11, 17, 21
- ▶ Outreach for OpenJDK 8 and 17 ELS opportunity
- ▶ OpenJDK Slim Images
 - Soon a Tech Preview.
- Based on UBI9 JDK21

JBoss Web Server

What's New in JBoss Web Server 6.0.x

Here are the highlights of the latest version:

- ▶ OpenJDK 21 Support
- ▶ Several CVEs fixed after the initial release

Please refer to the [release notes](#) for JBoss Web Server 6.0.x

Coming soon...

Some of the key features included in our roadmap for the next 12 months:

- ▶ **Runtimes Insights Integration**
- ▶ **UBI 9 Support**
- ▶ **Windows 2022 Support**
- ▶ **RHEL 10**
- ▶ **JWS 5 ELS-1 in starts in Aug 2025**
- ▶ And more...

Nodejs

What's in NodeJS

Here are the highlights of the latest offerings:

- ▶ Nodejs 18
- ▶ Nodejs 20
- ▶ Nodejs 22
- ▶ UBI 8 and UBI 9 support
- ▶ Key Modules Tested and Verified
 - Current [list](#)
- ▶ AI [Tutorials](#) for Nodejs

Please refer to the [release notes](#) for latest offerings

Coming soon...

Some of the key features included in our roadmap for the next 12 months:

- ▶ **Nodejs 24**
- ▶ **New blog posts about starting w/ Llama Stack**
- ▶ **Frequently evaluating new modules to be added**
- ▶ And more...



JBoss EAP 8



What's Next



★ What's New

► JBoss EAP 8.1 **Beta** Release

- Release Article on Red Hat Developers web site:
<https://developers.redhat.com/articles/2025/04/28/jboss-eap-8-1-beta-modernizing-enterprise-java>
- Continued Jakarta EE 10 compatibility for stability
- Java SE 17/21, Java SE 11 pruned (deprecated in EAP 7.4)
- Migration via Red Hat MTA product, Server Mig. Tool
- Bootable jar for JBoss EAP
- Server Footprint Trimming (using Galleon Technology) is on bare metal or VM environments
- Transition to VSCode for IDE Support

JBoss EAP 8.1 Beta: Modernizing enterprise Java applications

April 28, 2025 | Jeff Beck

Related topics: [Java](#), [Application modernization](#)

Related products: [Red Hat JBoss Enterprise Application Platform](#)

Share:

► JBoss EAP 8.1 **GA** Release

- Continued Jakarta EE 10 compatibility for stability
- Java SE 17/21, Java SE 11 pruned (deprecated in EAP 7.4)
- Migration via Red Hat MTA product, Server Mig. Tool
- Bootable jar for JBoss EAP
- Server Footprint Trimming (using Galleon Technology) is on bare metal or VM environments
- Transition to VSCode for IDE Support
- HashiCorp Credential Store Support
- Tested Configurations updates such as Oracle 23 AI, Postgresql 17, etc.
- Common Criteria Certification

► JBoss EAP XP 6 GA Release

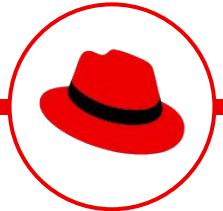
- MicroProfile 7 (OpenAPI 4.0, MicroProfile Telemetry 2.0, RESTClient 4.0, Fault Tolerance 4.1)
- Full Support for MicroProfile LRA
- Integrate OpenTelemetry with the MP Reactive messaging AMQP and Kafka Connectors.





streams for Apache Kafka

Optimization



1H '25

What's New

[Release Notes](#) for streams for Apache Kafka 2.9.0

- ▶ Last version compatible with Kafka clusters using ZooKeeper
- ▶ KRaft Generally Available
- ▶ End of Maintenance support-March 2026 and End of Limited Support - September 2026
- ▶ UI Console deployment via Operator GA
- ▶ Scale up/down rebalancing using CruiseControl
- ▶ Apache FlinkSQL demo improvements
- ▶ Please refer to the [release notes](#) for AMQ streams 2.9.0

Encryption

Stream processing

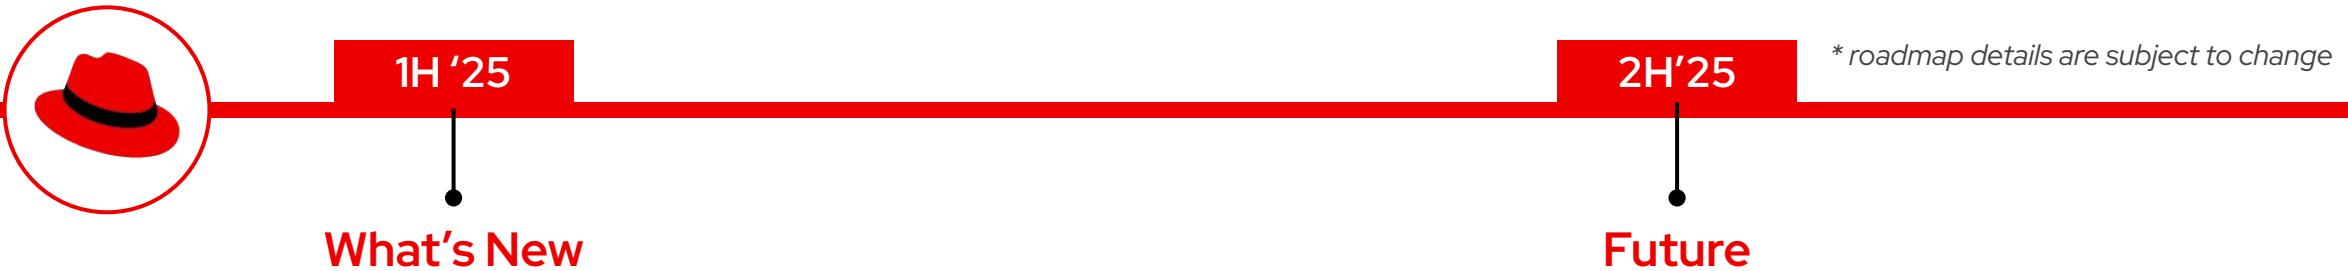
2H'25

* roadmap details are subject to change

Future

- ▶ KRaft Only
- ▶ KRaft: Scaling controller nodes
- ▶ Strimzi metrics reporter integration
- ▶ Cruise control self healing
- ▶ Tiered Storage support
- ▶ Proxy Operator
- ▶ Proxy: Record Encryption GA
- ▶ UI Console: Kafka Connect Tech preview
- ▶ Stream processing- FlinkSQL Tech preview
- ▶ Please refer to our [roadmap deck](#).

AMQ Broker



Release Notes for AMQ Broker 7.13 LTS

- ▶ Runs in FIPS-enabled OpenShift
- ▶ New AMQ Management Console based on HawtIO 4
 - Uses React and Patternfly
 - Faster and more responsive
 - Support for OAuth
- ▶ New percentage - based global memory limit setting
- ▶ New metrics exposed by the AMQ Prometheus plugin
- ▶ New Consumer priority support for AMQP Clients
- ▶ Federation stability improvements
- ▶ Improvements in JDBC performance on OpenShift

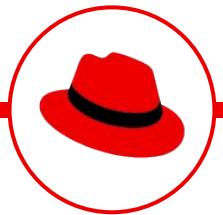
AMQ Broker 8.0 Tech Preview

- ▶ AMQ Operator and container images released independently of the Broker
- The Red Hat build of ArkMQ will replace AMQ Operator
- Faster releases decoupled from broker versions
- ▶ New CRDs
- Broker Provisioning CR creates a Broker Service
- Application CRs create application specific configurations for destinations, QoS, and authentication schemes etc.
- ▶ AMQP Bridges
- Aligns Federation, Mirroring, and Bridging all to AMQP protocol
- ▶ Improved Federation replaces Clustering as the default message distribution solution.



Red Hat 3scale API Management

General Enhancements



Q4 '24 - Q1'25

What's New

Security

Product Maintenance

Q3'25+

* roadmap details are subject to change

Highlights of [3scale v.2.15](#):

- ▶ Support for Red Hat build of Keycloak
- ▶ Support for Redis 7.2 as an external datastore
- ▶ Security, performance, and **general enhancements**, including:
 - Horizontal Pod Autoscaling across various components
 - APIcast compatibility with ARM64 architectures
 - Extended capabilities in 3scale Operator
 - Introduction of Openshift APIs for Data Protection as the *standard backup and restore mechanism*

Micro Releases [2.15.1](#), [2.15.2](#) and [2.15.3](#):

- ▶ Support for **ROSA HCP** (*Red Hat OpenShift Service on AWS with hosted control planes*)
- ▶ Support for **RHBK 26**
- ▶ Support for **OpenShift 4.18**
- ▶ Bug fixes and security enhancements

Red Hat 3scale API Management version [2.16](#):

- ▶ Support PostgreSQL 14
- ▶ Compliance with **Financial-grade API (FAPI) standards**
- ▶ TLS and ACL support for Redis connection
- ▶ Improved API endpoints for subscriptions and authentication options.
- ▶ Full support for the **Conditional Policy** in Apicast
- ▶ Option to disable integration with the **Zync component**
- ▶ Enhanced Audit Logs

Primary areas of focus for the future:

- ▶ Selective **RFE prioritization**.
- ▶ Ensuring cost-effective **product maintenance**
- ▶ Continued focus on **security** and addressing any reported vulnerabilities

Red Hat Build of Keycloak

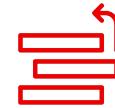
★ What's New

RHBK 26.2 (GA: May 13th, 2025) ; RH-SSO ELS-1 (2 years)

Key Highlights of the new version:

- ▶ Fine-Grained Admin Permissions (V2)
- ▶ Standard Token Exchange support (V2, as per rfc8693)
- ▶ ARM/Aarch64 ; ROSA clusters running on ARM
- ▶ Zero-configuration secure cluster communication
- ▶ Operator enhancements
- ▶ Observability enhancements
- ▶ Authentication improvements
- ▶ Improved Lifecycle Support Policies (see [Notes](#))

RH-SSO ELS-1 Availability (Jun 2025 - 2027) - more in [KB Article](#)



Coming soon...

RHBK 26.4 (Target GA: Q4 2025)

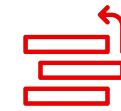
Key Highlights or focus areas for the upcoming version:

- ▶ Cross-Domain Token Exchange
- ▶ Automatic deletion and disabling of users (GDPR compliance)
- ▶ Admin Client v2 APIs
- ▶ Operator CRs for Clients
- ▶ Zero-downtime upgrades between patch releases for single site
- ▶ RHBK on a stretched OpenShift cluster across multiple sites
- ▶ Compatibility matrix (EDB PostgreSQL ; Azure SQL)
- ▶ Tech Preview features to GA fully supported
 - Passkeys ; Recovery Codes
 - Update Email Action ; DPoP

Red Hat build of Quarkus

★ What's New

- ▶ **Red Hat build of Quarkus 3.20 (May 2025)**
 - Enhanced observability with OpenTelemetry Logging
 - Full support of the new WebSocket implementation
 - Jakarta Websocket will be deprecated in this release, but we will keep it supported for at least RHBQ 3.x life cycle
 - Full support for generation of reflection-free Jackson serializers
 - Advanced Security Support for OIDC mTLS
 - Switch to defaulting to UBI9 Runtime images
 - Removal of RHBQ entitlement from RHEL 10



Coming soon...

- ▶ **Red Hat build of Quarkus 3.27 (ETA Oct 2025)**
 - ELS offering (or similar) extending life-cycle from 3 years to 5+ years
 - Launch of IBM build of Quarkus
 - Langchain4j support
 - MCP support
 - Including ELS
 - Subscription aware tooling (CLI, code.quarkus, etc)

Thank you

 linkedin.com/company/red-hat

 [youtube.com/user/RedHatVideo
s](https://youtube.com/user/RedHatVideos)

 facebook.com/redhatinc

 twitter.com/RedHat